

PyTorch Foundation

Technical Advisory Council Meeting

2026 September 15

Reminder for meeting participants

- This meeting is being recorded
- Please remember to be on mute upon entry unless you would like to speak
- All zoom chat messages will become public
- Please ensure your name and affiliation are showing correctly

Questions? Email support@pytorch.org

Antitrust Policy Notice

Linux Foundation meetings involve participation by industry competitors, and it is the intention of the Linux Foundation to conduct all of its activities in accordance with applicable antitrust and competition laws. It is therefore extremely important that attendees adhere to meeting agendas, and be aware of, and not participate in, any activities that are prohibited under applicable US state, federal or foreign antitrust and competition laws.

Examples of types of actions that are prohibited at Linux Foundation meetings and in connection with Linux Foundation activities are described in the Linux Foundation Antitrust Policy available at <http://www.linuxfoundation.org/antitrust-policy>. If you have questions about these matters, please contact your company counsel, or if you are a member of the Linux Foundation, feel free to contact Andrew Updegrove of the firm of Gesmer Updegrove LLP, which provides legal counsel to the Linux Foundation.



Agenda

- Administrative – 5 min
- Foundation Updates – 15 min
- Working Group Updates – 10 min
- Call for Topics/Open Discussion – 5 min

🕒 Meeting Schedule

2026

- ~~January 13~~
- ~~February 10~~
- ~~March 10~~
- ~~April 14~~
- ~~May 12~~
- ~~June 9~~
- ~~July 14~~
- ~~August 11~~
- September 15
- October 13
- November 10
- December 8

All meetings start at 7am PT/4pm CET unless otherwise indicated.

Upcoming TAC Meetups

PyTorch Conference North America

TAC Meetup (Members Only) - **RSVP to Calendar Invite by October 5th!**

Wednesday October 21 | 12:45pm - 2:15pm PT

[Register to Attend
Event](#)



Spread the word!



Share on your favorite social channels & tag us

- [LinkedIn](#)
- [X](#)
- [Facebook](#)

+ view our [Partner Marketing Kit](#) for suggested copy & graphics

Film a video!

- Find a quiet place, introduce yourself, detail how you'll be participating or what you're looking forward to at PyTorchCon NA, invite others to join you
- Send to marketing@pytorch.org + any social profiles you'd like tagged
- We'll post it to the PyTorch YouTube and social media channels

TAC Representatives – Member Companies

Welcome Davide
and Jing Zhu!

Company/Project	Primary Voting Representative	Alternate Voting Representative
Alibaba Cloud	Tao Ma	Sanhong Li
Amazon Web Services	Shauheen Zahirazami	Sundar Ranganathan
AMD	Jeff Daily	Niles Burbank
ARM	Milos Puzovic (TAC Vice Chair)	Andrew Wafaa
Cambricon	Jing Zhu	
Google Cloud	Claudio Basile	Michael Voznesensky
Huawei	Yikun Jiang	Fengchun Hua
IBM	Joseph Groenenboom	Naigang Wang
Intel	Eikan Wang	Ashok Emani
Lightning AI	Thomas Viehmann (TAC Chair)	
Meta	Gregory Chanan	Davide Italiano
NVIDIA	Piotr Bialecki	Michael Ruberry
Snowflake	Mert Hidayetoglu	
Shopify	Shuying Sun	
Gold Members	Weizhu Chen	

TAC Representatives – Projects

Company/Project	Primary Voting Representative	Alternate Voting Representative
PyTorch Project	Alban Desmaison	
DeepSpeed	Tunji Ruwase	Minjia Zhang
vLLM	Simon Mo	Kaichao You
Ray	Robert Nishihara	
Helion	Jongsok Choi	
Safetensors	Luc Georges	Lysandre Debut



Foundation Updates

Mark Collier

Time allocation: 15 minutes

New Members!

Platinum Members



Gold Members



New TAC Representatives



Jing Zhu

Cambricon[®]
寒 武 纪

🔄 Working Group Process and Lifecycle

- Current language on GitHub doesn't define process or lifecycle.
- Draft Working Group Formation Process and Lifecycle - link
 - Defines process for forming a working group, creating a charter, and regular reviews by the TAC.
- CTA: TAC to review and provide feedback



Working Group Updates

Time allocation: 10 minutes



CI Working Group Update

Thanh Ha


CI Working Group Update

- AWS Costs Update
 - Costs in August: 516k
 - AWS Credits remaining: 1.03m
 - 15% of jobs allocated to Foundation account (Excluding Inductor Jobs)
- GitHub OSDC (ARC) CI Updates
 - Continuing to optimize costs and improve cluster sizing
 - NAT Gateway fix to route S3 Traffic through S3 Gateway has had significant impact on transfer costs
 - Deployed changes to better spread pod load across all available base nodes allowing is to reduce base node size further.
- Working on new Cost Estimate & Runtime Reports webpage to replace the spreadsheets

Costs To PyTorch Foundation by Architecture

August 2026


Architecture	Cost
CUDA	\$194,725.11
x86_64 (AMD)	\$171,726.36
Other	\$104,477.12
ARM64 (AWS Graviton)	\$31,163.43
x86_64 (Intel)	\$13,867.13
Windows	\$1.59
s390x	\$0.00
TPU	\$0.00
x86_64 (Unknown)	\$0.00
XPU	\$0.00
ROCM	\$0.00
ARM64 (Unknown)	\$0.00
macOS	\$0.00
Total	\$515,960.74



* Data from PT Foundation AWS Account. This does not include self-hosted runner costs by member provided runners.

Combined Hardware Runtime by Architecture

12 month trend (1/2)



* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners.

Combined Hardware Runtime by Architecture

12 month trend (2/2)

Architecture	2025-09	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	2026-06	2026-07	2026-08
x86_64 (AMD)	17,924	18,324	13,524	17,403	15,138	16,026	37,035	173,961	891,472	1,185,121	940,734	1,381,359
CUDA	259,275	257,990	200,793	333,761	352,590	348,113	389,233	384,112	578,941	556,751	390,611	653,054
x86_64 (Intel)	765,559	759,300	589,811	688,334	638,775	672,990	978,942	917,098	780,824	402,353	308,600	534,842
ARM64 (AWS Graviton)	91,424	76,686	65,816	67,670	60,321	60,254	99,039	155,932	235,347	210,911	164,301	266,768
ROCM	77,547	84,025	60,722	132,535	84,198	115,055	147,636	182,744	207,576	212,748	224,600	217,024
macOS	47,606	48,311	30,197	34,581	32,957	38,545	53,188	60,221	64,507	53,691	55,018	170,877
Windows	118,229	84,008	61,171	68,830	56,358	58,356	83,315	107,058	112,621	88,879	101,692	149,575
s390x	6,363	2,374	1,978	2,090	2,246	3,182	2,657	8,172	6,009	2,366	3,554	2,899
TPU	0	0	0	73	126	454	891	1,701	1,270	855	928	901
XPU	93,246	83,311	66,666	106,968	96,457	80,503	103,778	124,526	166,931	95,992	112,785	136,591
x86_64 (Unknown)	98	108	107	114	51	67	678	583	666	1,050	1,538	2,397
ARM64 (Unknown)	0	0	0	0	0	7	0	0	0	0	0	450
Other	0	0	0	0	0	0	0	0	0	0	0	0
N/A	-2	0	0	0	0	-46	-35	-73	-71	0	0	0
Total	1,477,268	1,414,436	1,090,785	1,452,358	1,339,218	1,393,504	1,896,358	2,116,036	3,046,092	2,810,717	2,304,361	3,516,737

* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners.



Ecosystem Working Group

Chris Hoge/Eikan Wang

Ecosystem Working Group Updates

Accepted 🎉:

- 🎉 [TorchSurv](#) - "Deep survival analysis made easy"
- 🎉 [torchao](#) - easy to use quantization library for native PyTorch
- 🎉 [Zeus](#) - library for measuring the energy consumption of Deep Learning workloads

Active Applications 📦:

- 📦 [LiteRT](#) - new application for on-device inference - needs intake review

Under Review 🔍:

- ⚠️ [LitData](#) - Speed up model training by fixing data loading - *Addressing governance, CODEOWNERS and PyTorch version policy*
- ⚠️ [FlagGems](#) - A Triton-based Operator Library for PyTorch - *Governance addressed; CI fixes requested*
- ⚠️ [PyTorch-Wildlife](#) - A platform designed for creating, modifying, and sharing powerful AI models - *Awaiting a name change, with development paused*
- ⚠️ [NeuralDBG](#) - A causal diagnostic engine for PyTorch training - *Too early for inclusion, pending broader community participation*

Upcoming Workstreams 🛠️:

- Writing new WG charter
- Finalizing survey
- Evaluation skill merged
- Preparing talk for PyTorch Conference

Monthly meeting on last Wednesday
at 9am PT/12pm ET - [link](#)

<https://github.com/pytorch-fdn/ecosystem>
Slack: [#pytorch-ecosystem-working_group](#)



Accelerator Integration Working Group

Zesheng Zong

Accelerator Integration Working Group Updates

Cross-Repository CI Relay ([RFC](#))

- L3 promotion criteria – L3 promotion criteria is ready ([#102](#)).
- L3 readiness panel – Shipped the per-repo promotion/demotion panel with fixed 14d/7d windows, tenure tracking, and max execution / E2E metrics ([#8658](#), [#8635](#), [#8693](#), [#8707](#)).
- L3 bugs & Relay accuracy – Fixed HUD display and check-run creation for push events, workflow-level queue time tracking, max execution time calculation excluding timeout jobs, retry double-counting, push-event PR context recovery, and PR matrix reporting to only include the latest commit ([#8707](#), [#8666](#), [#8557](#), [#8555](#), [#8620](#), [#8730](#)).
- Nightly HUD – Shipped the Nightly tab, per-repo nightly view, and PR/Nightly metrics tabs; fixed summary/matrix consistency issues including partial rows, SHA-grouped stats, and hiding crcr-test ([#8377](#), [#8441](#), [#8538](#), [#8695](#), [#8696](#), [#8697](#)).
- RFC-0054 payload gaps – Added failed-test ingestion, drill-down UI, and artifact_url rendering on the per-repo dashboard ([#8607](#), [#8608](#), [#8618](#)).
- External CI / Buildkite – Added multi-issuer OIDC and ci_providers.yml, registered vLLM Buildkite, and enabled external nightly results through Results Relay ([#8453](#), [#8468](#), [#8544](#), [#8326](#)).
- Nightly hardening – Working on write-once protection for finalized nightly callbacks and refining RFC-0056; xfail/xcancel policy remains open ([#8676](#), [#8694](#), [#8623](#), [RFC #98](#)).

Test Refactoring

- **276/1,023** test files have been refactored, with **130+** PRs under initial review and **181** PRs under final review.



Security Working Group

Nikita Shulga

Security WG Updates

- 20 issues reported, 21 closed, 4 in draft, 0 published
 - Removed NestedTensors from safe list
 - Most of the valid issues are in weights_only feature
- Incoming advisories are triaged every other week on Tuesdays at 8:30AM PST
 - Meetings are not open to public by design
 - If you have some security background and want to join, post in `#tac-security-wg`



OSPO & Academic Outreach Working Group

Ana Jimenez/George Chellapa/Sumanthro Mukherjee

OSPO & Academic Outreach WG Updates

Repo: <https://github.com/pytorch-fdn/wg-ospo-and-academic-outreach>

WG Chairs: Sumantra Mukherjee (Red Hat), George Chellapa (NVIDIA), Ana Jiménez

Updates:

- Finalized a Day Zero Academic Community Event proposal for PyTorch Conference North America (#17)
- Call for Projects will be out soon along with a blogpost
- OSPO & Academic Outreach WG working with Events and Marketing team to announce it on the Website.
- Pilot blog highlighting TinyTorch being drafted in collaboration with Harvard University.

Open Discussion

Open Discussion

- Any member can take the floor

Next Meeting

Next regular TAC meeting

Tuesday, October 13, 2026

7am PDT / 10am EDT / 4pm CEST

Topics: tbd

Please ensure you are registered for the meeting series and that meetings are reflected in your calendar via [LFX Meetings](#)

Have an agenda topic idea? Post topics to the [TAC Agenda Github Project](#)

🔗 Helpful Links

Links

Foundation GitHub:

<https://github.com/pytorch-fdn/tac>

Mailing List:

<https://lists.pytorch.org/g/tac>

Past Meeting Minutes and Decks:

<https://github.com/pytorch-fdn/tac/tree/main/minutes>

LFX Individual Dashboard:

<https://openprofile.dev>

Submit an Ecosystem project:

<https://github.com/pytorch-fdn/ecosystem>

Submit a Blog:

<https://form.asana.com/?k=mPrsX9EcZ0BoGnqz543C1A&d=9283783873717>

Points of Contact

[Mark Collier](#) (Executive Director)

[Naomi Washington](#) (Senior Program Manager)

[Jennifer Bly](#) (Marketing Director)

[Regina Nkenchor](#) (Program Manager)

[Michelle Roth](#) (Program Manager)

Questions? Email support@pytorch.org



Appendix

🕒 Foundation-Hosted Project Review Schedule

- January 13, 2026
- April 14, 2026
- July 14, 2026
- November 10, 2026 – updated



Events



Check out the
[PyTorch Events Calendar](#)



Marketing

JOIN US

PyTorch Conference North America

October 20-21, 2026

San Jose, CA


#PyTorchCon

[SUBMIT TO SPEAK](#)[SPONSOR](#)

Get Started: Install PyTorch Locally or Launch Instantly on Supported Cloud Platforms

[Get started](#)

Visit the PyTorch website: <https://pytorch.org/>



PyTorch ✓

@PyTorch

Tensors and neural networks in Python with strc
PyTorch is an open source project at the Linux F

pytorch.org Joined September 2016 >

86 Following 496.5K Followers



What PyTorch Conference Europe 2026 Was
Really Like – Official PyTorchCon EU ...

PyTorch
321,544 followers
1w · Edited

Using PyTorch allowed [LinkedIn](#) to enable production-grade optimization at previously infeasible scales.


In our latest PyTorch Foundation case study, the engineering team at LinkedIn shares how they re-architected their extreme scale Linear Programming Solver (DuaLip GPU). Built using PyTorch to enable multi-GPU computations and parallelism, this solver serves as a core building block for massive matching problems that power ranking, personalization, and recommendation systems.

To maximize parallelism, the team tailored GPU execution techniques to sparse matching constraints. They developed constraint-aligned sparse layouts, batched projection kernels, and a distributed design focused on dual variables. They also updated the underlying ridge-regularized dual ascent method with Jacobi-style row normalization, primal scaling, and a regularization parameter continuation scheme.

When tested against extreme scale matching workloads, the GPU implementation achieved a 10x wall-clock speedup over the previous distributed CPU DuaLip solver while maintaining strict convergence guarantees. This project represents an excellent intersection of ML systems, mathematical optimization, and machine learning.

Read the complete case study here <https://lnkd.in/eE67sxcB>

Aida Rahmattalabi Sanjana Garg Zhipeng Wang, PhD Shenyinying Ruby Tu Yuan G.
Gregory Dexter



Follow PyTorch on social media: [X](#), [LinkedIn](#), [Facebook](#), [Bluesky](#), [YouTube](#)

Sign up today!



Subscribe to the PyTorch Newsletter to get the latest updates, event info, and community news in your inbox each month → pytorch.org/newsletter

PyTorch Trademark Policy

PyTorch® is a registered trademark. Per the Linux Foundation Trademark Usage Guidelines:

Trademarks are used to provide assurance to the community of consistency with respect to the quality of products or services with which the mark is associated. This policy attempts to strike the proper balance between two competing interests:

- 1) the need of The Linux Foundation to ensure that its trademarks remain reliable indicators of the qualities that they have been created to preserve; and,
- 2) the need of The Linux Foundation to ensure that community members are able to discuss the projects with which The Linux Foundation is associated and to accurately describe the relationship between The Linux Foundation and the products and services offered by others.

Questions? Email trademark@pytorch.org

- **Don't** use a trademark as a verb or noun - use only as an adjective followed by the generic name/noun.
- **Don't** use in the plural or possessive form.
- **Don't** alter or amend a trademark in any way, including combining with another mark, hyphenating, abbreviating, or displaying in parts.
- **Don't** use it as your domain or as part of your domain name.
- **Don't** use it as part of your product name.
- **Don't** incorporate it into your company's logos or designs.
- **Do** make fair use of word marks to make true factual statements, as long as you acknowledge the owner of the trademark.
- **Do** use the full form of any trademark in the first reference, after which you may then use any abbreviated or short form references.

For the full trademark guidance, please review

<https://www.linuxfoundation.org/legal/trademark-usage> and
https://pytorch.org/wp-content/uploads/2025/09/pytorch_brand_guide_091925a.pdf

Legal Notices

The Linux Foundation, The Linux Foundation logos, and other marks that may be used herein are owned by The Linux Foundation or its affiliated entities, and are subject to The Linux Foundation's Trademark Usage Policy at <https://www.linuxfoundation.org/trademark-usage>, as may be modified from time to time.

Linux is a registered trademark of Linus Torvalds. Please see the Linux Mark Institute's trademark usage page at <https://lmi.linuxfoundation.org> for details regarding use of this trademark.

Some marks that may be used herein are owned by projects operating as separately incorporated entities managed by The Linux Foundation, and have their own trademarks, policies and usage guidelines.

TWITTER, TWEET, RETWEET and the Twitter logo are trademarks of Twitter, Inc. or its affiliates.

Facebook and the "f" logo are trademarks of Facebook or its affiliates.

LinkedIn, the LinkedIn logo, the IN logo and InMail are registered trademarks or trademarks of LinkedIn Corporation and its affiliates in the United States and/or other countries.

YouTube and the YouTube icon are trademarks of YouTube or its affiliates.

All other trademarks are the property of their respective owners. Use of such marks herein does not represent affiliation with or authorization, sponsorship or approval by such owners unless otherwise expressly specified.

The Linux Foundation is subject to other policies, including without limitation its Privacy Policy at <https://www.linuxfoundation.org/privacy> and its Antitrust Policy at <https://www.linuxfoundation.org/antitrust-policy>, each as may be modified from time to time. More information about The Linux Foundation's policies is available at <https://www.linuxfoundation.org>.

Please email legal@linuxfoundation.org with any questions about The Linux Foundation's policies or the notices set forth on this slide.