PyTorch Foundation Technical Advisory Council Meeting

2026 October 13

Reminder for meeting participants

- This meeting is being recorded
- Please remember to be on mute upon entry unless you would like to speak
- All zoom chat messages will become public
- Please ensure your name and affiliation are showing correctly

Questions? Email support@pytorch.org

Antitrust Policy Notice

Linux Foundation meetings involve participation by industry competitors, and it is the intention of the Linux Foundation to conduct all of its activities in accordance with applicable antitrust and competition laws. It is therefore extremely important that attendees adhere to meeting agendas, and be aware of, and not participate in, any activities that are prohibited under applicable US state, federal or foreign antitrust and competition laws.

Examples of types of actions that are prohibited at Linux Foundation meetings and in connection with Linux Foundation activities are described in the Linux Foundation Antitrust Policy available at http://www.linuxfoundation.org/antitrust-policy. If you have questions about these matters, please contact your company counsel, or if you are a member of the Linux Foundation, feel free to contact Andrew Updegrove of the firm of Gesmer Updegrove LLP, which provides legal counsel to the Linux Foundation.

Agenda

- Administrative - 5 min
- Foundation Updates - 30 min
- Working Group Updates - 15 min
- Call for Topics/Open Discussion - 5 min

Meeting Schedule

2026

- ~~January 13~~
- ~~February 10~~
- ~~March 10~~
- ~~April 14~~
- ~~May 12~~
- ~~June 9~~
- ~~July 14~~
- ~~August 11~~
- ~~September 15~~
- October 13
- November 10
- December 8

All meetings start at 7am PT/4pm CET unless otherwise indicated.

Upcoming TAC Meetups

PyTorch Conference North America

TAC Meetup (TAC and GB Members Only)

Wednesday October 21 | 12:45pm - 2:15pm PT

Register to Attend Event



Spread the word!



Share on your favorite social channels & tag us

- LinkedIn
- X
- Facebook

+ view our Partner Marketing Kit for suggested copy & graphics

Film a video!

- Find a quiet place, introduce yourself, detail how you'll be participating or what you're looking forward to at PyTorchCon NA, invite others to join you
- Send to marketing@pytorch.org + any social profiles you'd like tagged
- We'll post it to the PyTorch YouTube and social media channels

Spread the word!

- Submit your event by October 15!
- https://events.linuxfoundation.org/open-source-ai-week/



Film a Short Video with Us!

We’re filming short video interviews onsite at PyTorch Conference next week spotlighting our members, projects, and community.

Topic suggestions:

- Why you’re are a part of the TAC
- How you use PTF projects
- Your PyTorch Conference experience

Interested in filming a quick spot?

Book a time (calendly link) or reach out to marketing@pytorch.org and we’ll get you set up.



TAC Representatives - Member Companies

Company/Project	Primary Voting Representative	Alternate Voting Representative
Alibaba Cloud	Tao Ma	Sanhong Li
Amazon Web Services	Shauheen Zahirazami	Sundar Ranganathan
AMD	Jeff Daily	Niles Burbank
ARM	Milos Puzovic (TAC Vice Chair)	Andrew Wafaa
Cambricon	Jing Zhu	
Google Cloud	Claudio Basile	Michael Voznesensky
Huawei	Yikun Jiang	Fengchun Hua
IBM	Joseph Groenenboom	Naigang Wang
Intel	Eikan Wang	Ashok Emani
Lightning AI	Thomas Viehmann (TAC Chair)	
Meta	Gregory Chanan	Davide Italiano
NVIDIA	Piotr Bialecki	Michael Ruberry
Snowflake	Mert Hidayetoglu	
Shopify	Shuying Sun	
Gold Members	Weizhu Chen	

TAC Representatives - Projects

Company/Project	Primary Voting Representative	Alternate Voting Representative
PyTorch Project	Alban Desmaison	
DeepSpeed	Tunji Ruwase	Minjia Zhang
vLLM	Simon Mo	Kaichao You
Ray	Robert Nishihara	
Helion	Jongsok Choi	
Safetensors	Luc Georges	Lysandre Debut

Foundation Updates
Mark Collier

Time allocation: 30 minutes

Foundation-Hosted Projects

The PyTorch Foundation is the vendor-neutral home for the open source intelligence layer developers use for training, optimizing, serving, orchestrating, and running models on any chip in any cloud for any agent.

As a community-driven hub hosted by the Linux Foundation, the PyTorch Foundation supports the core PyTorch framework alongside a growing portfolio of innovative projects including vLLM, DeepSpeed, Ray, Helion, and Safetensors.

Through open governance, strategic support, and a global contributor community, the PyTorch Foundation empowers developers, researchers, and enterprises to build and deploy AI at scale.

- Discussion: What projects would fit into this mission/align with Foundation goals?

Foundation-Hosted Project Lifecycle Process Update

Latest version here: [link]

5 Improvements Over Lifecycle v1.0:

1	**Clearer maturity ladder**	Incubated, Graduated and Archived replace one hosted stage, Candidate status and Platform/Vertical types.
2	**Simpler on-ramp**	No Ecosystem prerequisite or Candidate step: submit a proposal, present to the TAC, win a majority vote.
3	**A reason to keep growing**	Graduation (TAC and Governing Board supermajority) adds a TAC seat and priority event access.
4	**More ongoing accountability**	Quarterly updates, annual TAC review, public quarterly TSC meetings, EasyCLA or DCO.
5	**Shorter, more usable doc**	Judgment-based criteria, a separate guidance doc, and clearer archive steps (README status, successors).

Working Group Updates

Time allocation: 15 minutes

Working Group Process and Lifecycle

- Current language on GitHub doesn't define process or lifecycle.

- Draft Working Group Formation Process and Lifecycle - link
  - Defines process for forming a working group, creating a charter, and regular reviews by the TAC.

- CTA: TAC to vote to approve



CI Working Group Update

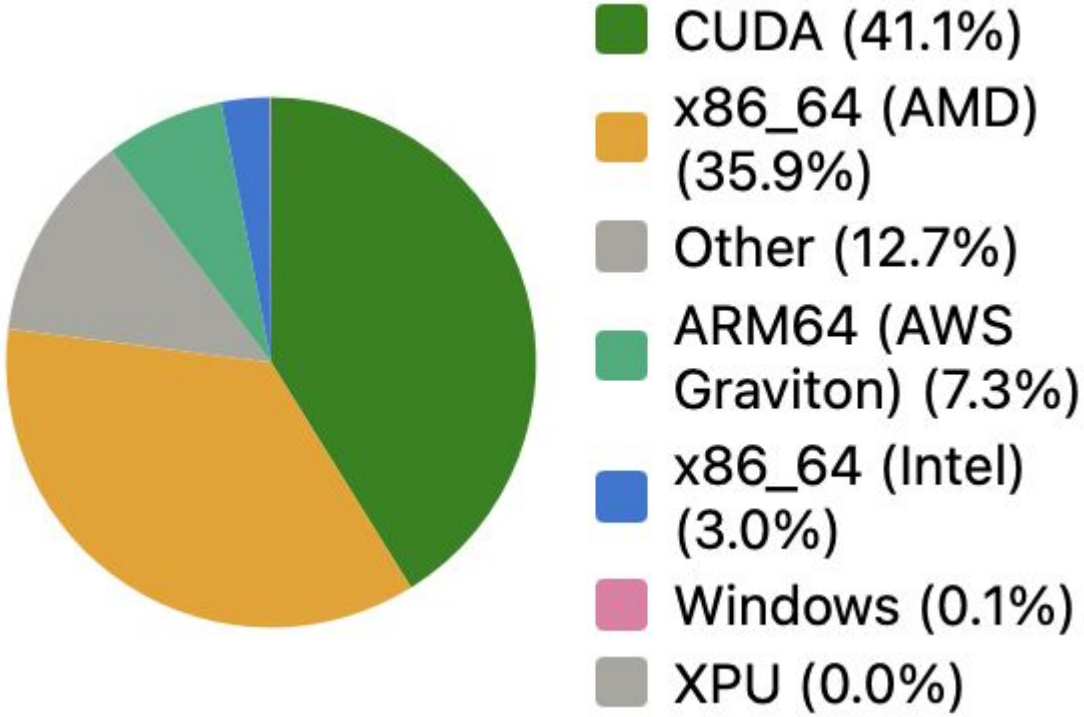
Thanh Ha

CI Working Group Update

- AWS Costs Update
  - Costs in September: 452k
  - AWS Credits remaining: 1.07m
  - 20% of jobs allocated to Foundation account (Excluding Inductor Jobs)
- GitHub OSDC (ARC) CI Updates
  - Continuing to optimize costs and improve cluster sizing
  - Working on runner allow list to give us more flexibility in which runners run on the LF Fleet
  - Tuned G4DN, G5, and G6 instance type priority on LF Fleet to reduce overhead costs
- New Cost Estimate & Runtime Reports webpage replaces the spreadsheets
  - Available at https://ci-reports.pytorch.org/
    - Requires LFID: Reach out to Thanh to get access

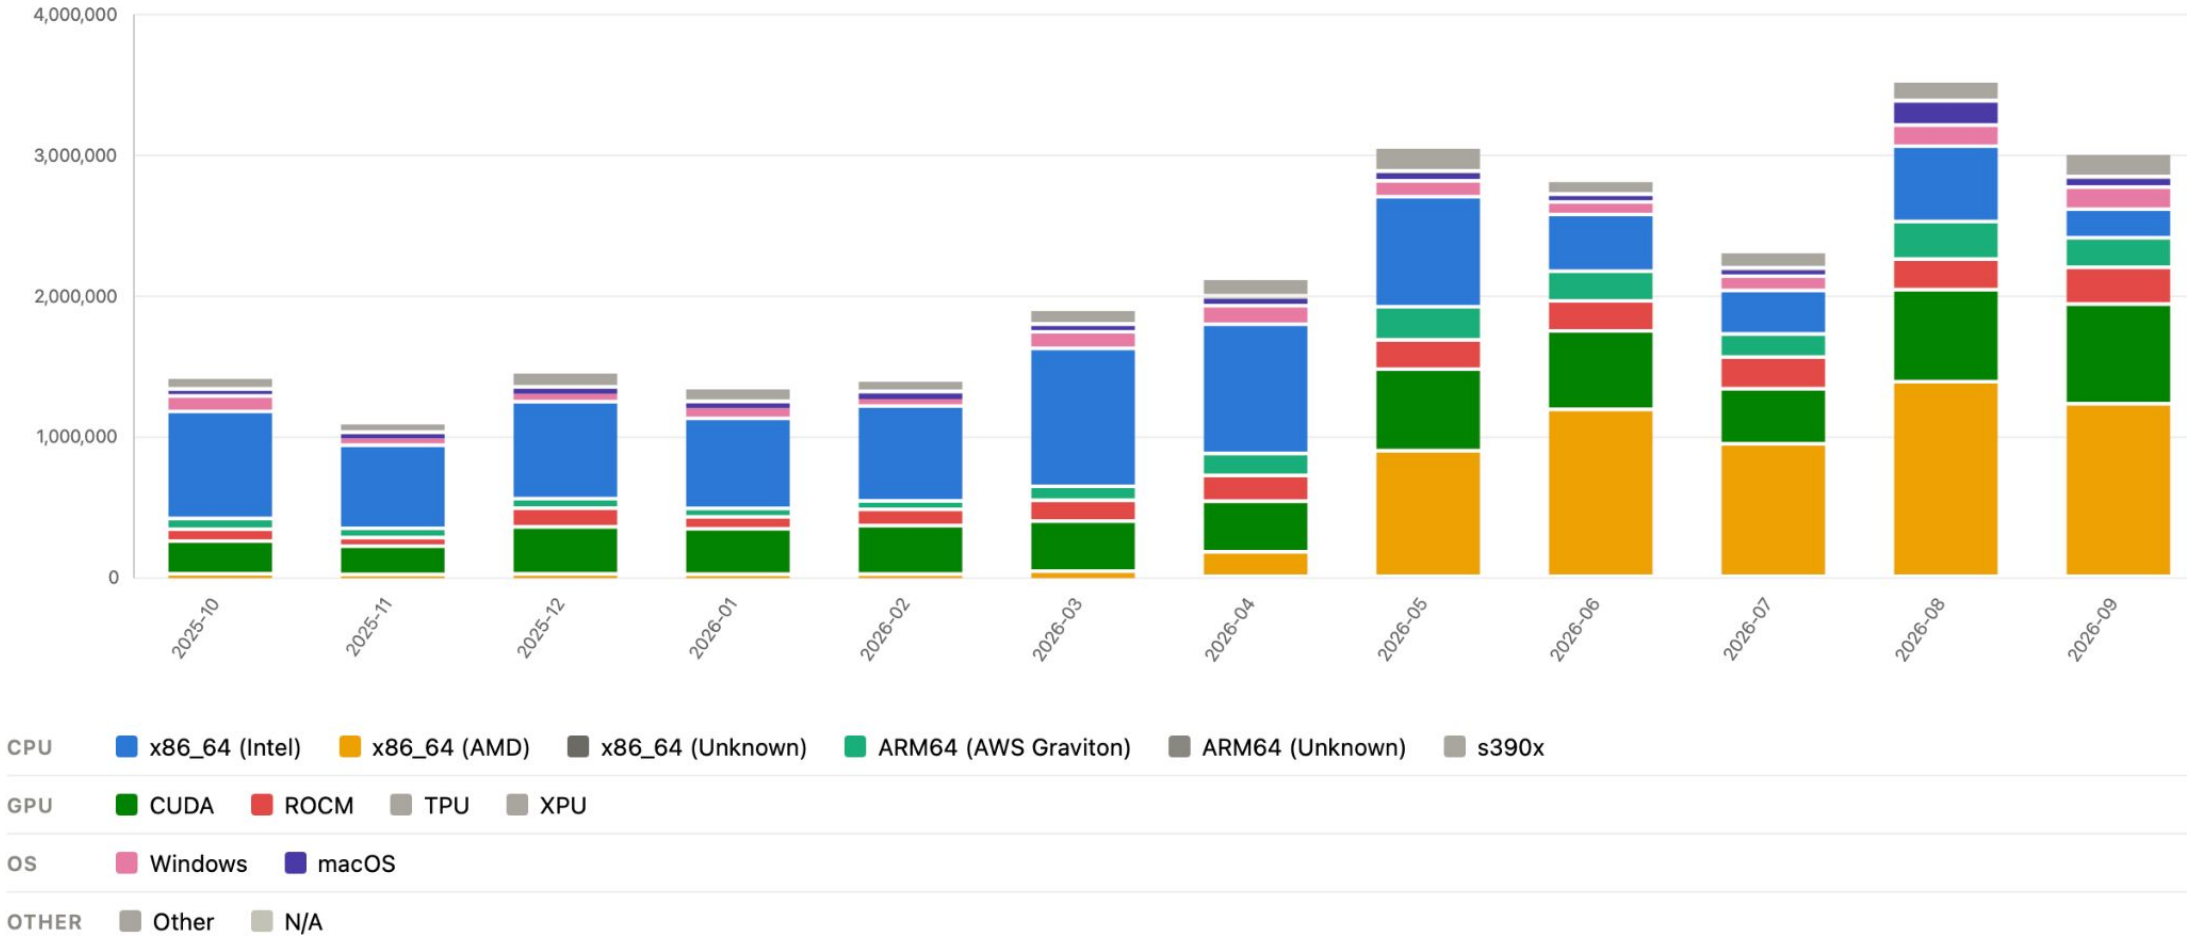
Costs To PyTorch Foundation by Architecture September 2026

Architecture	Cost
CUDA	$185,781.67
x86_64 (AMD)	$162,076.68
Other	$57,233.66
ARM64 (AWS Graviton)	$32,770.04
x86_64 (Intel)	$13,385.46
Windows	$276.75
XPU	$16.82
TPU	$0.00
ROCM	$0.00
s390x	$0.00
ARM64 (Unknown)	$0.00
macOS	$0.00
x86_64 (Unknown)	$0.00
N/A	$-16.82
Total	**$451,524.26**



* Data from PT Foundation AWS Account. This does not include self-hosted runner costs by member provided runners. – https://ci-reports.pytorch.org/

Combined Hardware Runtime by Architecture
12 month trend (1/2)



* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners. – https://ci-reports.pytorch.org/

Combined Hardware Runtime by Architecture 12 month trend (2/2)

Architecture	Location	Funded By	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	2026-06	2026-07	2026-08	2026-09
x86_64 (AMD)	AWS, GitHub-hosted (Azure)	Amazon, Meta	18,324	13,524	17,403	15,138	16,026	37,035	173,961	891,472	1,185,121	940,734	1,381,359	1,224,913
CUDA	AWS, GCP, Nvidia lab, Unknown	Amazon, Meta, Nvidia, Unknown	231,857	200,793	333,761	322,462	344,206	355,088	359,527	578,941	556,751	390,611	653,054	707,582
ROCM	AMD lab, DigitalOcean, Vultr	AMD	84,025	60,722	132,535	84,198	115,055	147,636	182,744	207,576	212,748	224,600	217,024	260,580
ARM64 (AWS Graviton)	AWS	Amazon, Meta	76,686	65,816	67,670	60,321	60,254	99,039	155,932	235,347	210,911	164,301	266,768	210,066
x86_64 (Intel)	AWS	Amazon, Meta	759,300	589,811	688,334	638,775	672,990	978,942	917,098	780,824	402,353	308,600	534,842	202,687
Windows	AWS, GitHub-hosted (Azure)	Amazon, Meta	110,140	61,171	68,830	86,486	62,262	117,425	131,555	112,621	88,879	101,692	149,575	156,548
macOS	GitHub-hosted (Apple silicon), Unknown	Meta, Unknown	48,311	30,197	34,581	32,957	38,545	53,188	60,221	64,507	53,691	55,018	170,877	69,505
s390x	IBM	IBM	2,374	1,978	2,090	2,246	3,182	2,657	8,172	6,009	2,366	3,554	2,899	5,027
TPU	Google Cloud	Google	0	0	73	126	454	891	1,701	1,270	855	928	901	1,122
XPU	AWS, Intel lab	Amazon, Meta, Intel	83,311	66,666	106,968	96,457	80,503	103,778	124,526	166,931	95,992	112,785	136,591	160,280
x86_64 (Unknown)	GitHub-hosted (Azure)	Meta	108	107	114	51	67	678	583	666	1,050	1,538	2,397	2,611
ARM64 (Unknown)	GitHub-hosted (Azure)	Meta	0	0	0	0	7	0	0	0	0	0	450	1,169
Other	AWS	Amazon	0	0	0	0	0	0	0	0	0	0	0	0
N/A	Unknown	Unknown	0	0	0	0	-46	0	15	-71	0	0	0	0
Total			**1,414,436**	**1,090,785**	**1,452,358**	**1,339,218**	**1,393,504**	**1,896,358**	**2,116,036**	**3,046,092**	**2,810,717**	**2,304,361**	**3,516,737**	**3,002,091**

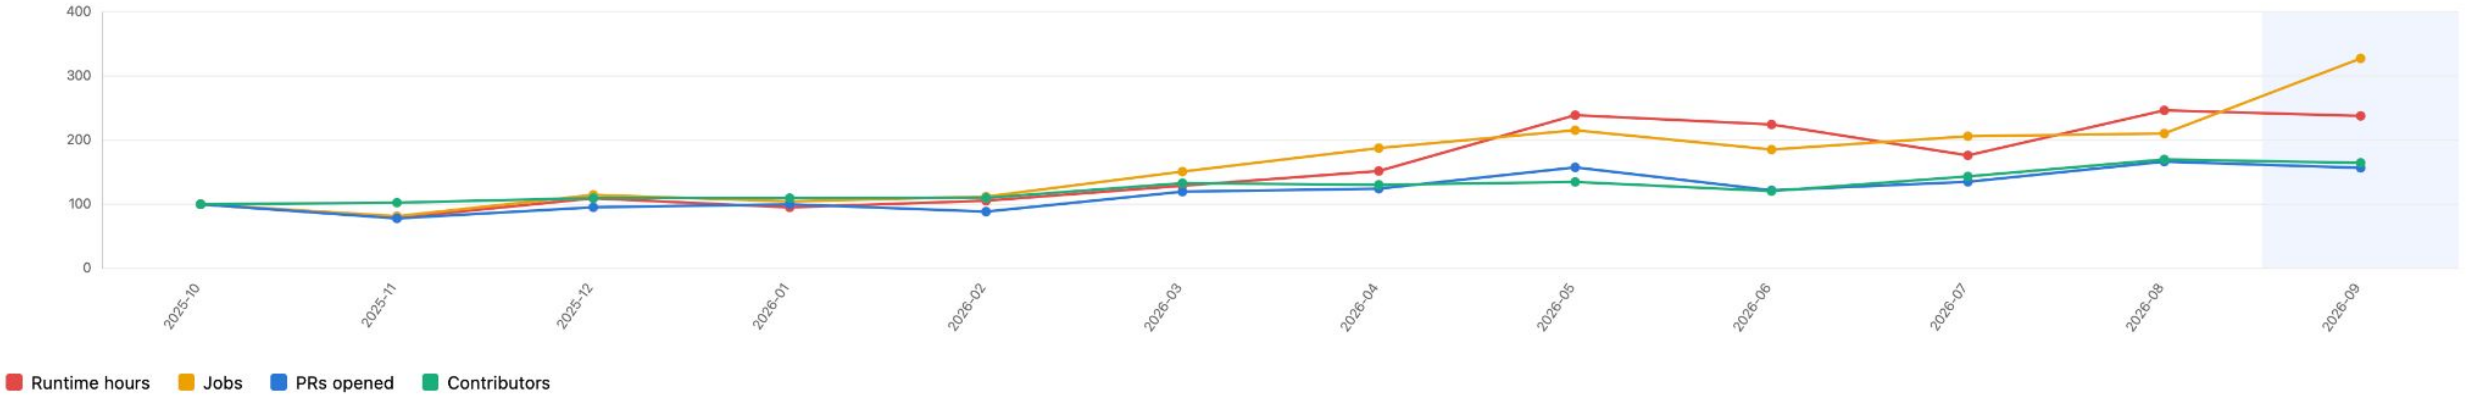
* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners. – https://ci-reports.pytorch.org/

CI Growth Metrics - 12 month trend

Repos: pytorch/pytorch only Month: 2026-09 pytorch/pytorch only.

PRs opened	Jobs run	Contributors	Runtime hours	Hours / PR	Failed + cancelled hours
3K	**5.5M**	**591**	**2.4M**	**819.9**	**16.2%**
▼6% MoM ▲66% YoY	▲56% MoM ▲214% YoY	▼3% MoM ▲68% YoY	▼4% MoM ▲132% YoY	▲2% MoM ▲40% YoY	▼17.0 pts MoM ▲7.1 pts YoY

Growth vs. first month with all data in window (index = 100)



* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners. – https://ci-reports.pytorch.org/

Ecosystem Working Group

Chris Hoge/Eikan Wang

Ecosystem Working Group Updates

- Working Group Charter
  - How do we position the Ecosystem Working group in relation to PyTorch Foundation hosted projects in a meaningful way that is not misleading to the community?
  - ***"The PyTorch Ecosystem strengthens the PyTorch Landscape by recognizing open, community-oriented projects outside the Foundation. Participation provides a low-barrier engagement that supports the growth and visibility of independent projects that actively contribute to the health of the PyTorch community."***
  - The Ecosystem can remain a stepping stone to full PyTorch Foundation hosted status, but it can also be an endpoint for projects that are committed to the community.
  - Looking for guidance from Foundation leadership on the best positioning and messaging.
  - Ecosystem project survey is running right now.

Monthly meeting on last Wednesday
at 8am PT/12pm ET -
https://github.com/pytorch-fdn/ecosystem
Slack: #pytorch-ecosystem-working_group

Ecosystem Working Group Updates

Accepted 🎉:

- LitData - *provides tools to preprocess and optimize datasets into a format that streams efficiently from any cloud or local source. It also includes a map operator for distributed data processing before optimization.*
- XGrammar - *a widely-adopted open-source library for efficient structured generation in PyTorch-based large language model inference.*

Active Applications 🗳️:

- LiteRT - out for voting - *LiteRT is Google's high-performance, open-source runtime and conversion toolchain for deploying PyTorch models on edge, mobile, desktop, web, and embedded devices.*

Under Review 🔍:

- *None under active review*

Renewed 🛠️:

- torchdistill - *offers various state-of-the-art knowledge distillation methods and enables you to design (new) experiments simply by editing a declarative yaml config file instead of Python code.*

Accelerator Integration Working Group

Zesheng Zong

Accelerator Integration Working Group Updates

Cross-Repository CI Relay (RFC)

- L3 gating became operational — Readiness now uses overrun rate and L2 tenure. An open demotion PR, which a daily sync can create, suppresses new upstream check runs without stopping downstream observation. Downstreams can also attach a triage verdict to callbacks to tell upstream-induced failures from backend issues; HUD rendering of the verdict is in review. (#8754, #8792, #8857, #8941, #8944, #8803, #9018)
- Event-scoped backend participation shipped — Backends can opt into pull_request, nightly, or both. Nightly-only repos are excluded from PR views, dispatch, and L3 backfill. (#8854, #8855, #8856, #8862)
- Dr.CI shows L3 results only when the CRCR label is present — Unlabeled shadow work no longer appears as a PR check. (#8881)
- Nightly data and health diagnosis hardened — Fixes cover callback ingestion, Buildkite mappings, SHA-based grouping, latest-attempt counts, and UTC timestamps. Finalized results are now write-once, overdue probes degrade health, and health cards link to problem jobs. (#8843, #8836, #8860, #8832, #8896, #8694, #8801, #8834)
- Open items — Nightly pass-rate aggregation was reverted and is back in draft; it awaits ClickHouse triage with the related summary and classification queries. The callback-rate-limit incident is still open, and per-repo level history is in review. (#8927, #9024, #9061, #8961, #9045, #9046, #9019, #8894)

Test Refactoring

- 397/1,023 test files have been refactored, with 108 PRs under initial review and 14 PRs under final review.

Planning Working Group In-Person Meeting @ PyTorch Conference NA 2026



Security Working Group

Nikita Shulga

Security WG Updates

to be updated

- 20 issues reported, 21 closed, 4 in draft, 0 published
  - Removed NestedTensors from safe list
  - Most of the valid issues are in weights_only feature

- Incoming advisories are triaged every other week on Tuesdays at 8:30AM PST
  - Meetings are not open to public by design
  - If you have some security background and want to join, post in `#tac-security-wg`

OSPO & Academic Outreach Working Group

Ana Jimenez/George Chellapa/Sumantro Mukherjee

OSPO & Academic Outreach WG Updates

Repo: https://github.com/pytorch-fdn/wg-ospo-and-academic-outreach

WG Chairs: Sumantro Mukherjee (Red Hat), George Chellapa (NVIDIA), Ana Jiménez

Updates:

- Day Zero Academic Community Event confirmed for October 19th 9:00 A.M. to noon (PST).
- Registration Open.
- Still looking for volunteers to help with the workshop - how to move a project from research to a sustainable open source project.
- Pilot blog highlighting TinyTorch live!
- We welcome ideas for future blog posts highlighting academic projects.

Open Discussion

Open Discussion

- Any member can take the floor

Next Meeting

Next regular TAC meeting

Tuesday, November 10, 2026

7am PDT / 10am EDT / 5pm CET

Topics: Q4 Foundation-Hosted Project Updates

Please ensure you are registered for the meeting series and that meetings are reflected in your calendar via LFX Meetings

Have an agenda topic idea? Post topics to the TAC Agenda Github Project

Helpful Links

Links

Foundation GitHub:
https://github.com/pytorch-fdn/tac

Mailing List:
https://lists.pytorch.org/g/tac

Past Meeting Minutes and Decks:
https://github.com/pytorch-fdn/tac/tree/main/minutes

LFX Individual Dashboard:
https://openprofile.dev

Submit an Ecosystem project:
https://github.com/pytorch-fdn/ecosystem

Submit a Blog:
https://form.asana.com/?k=mPrsX9EcZ0BoGnqz543C1A&d=9283783873717

Points of Contact

Mark Collier (Executive Director)

Naomi Washington (Senior Program Manager)

Jennifer Bly (Marketing Director)

Regina Nkenchor (Program Manager)

Michelle Roth (Program Manager)

Questions? Email support@pytorch.org

Appendix

Foundation-Hosted Project Review Schedule

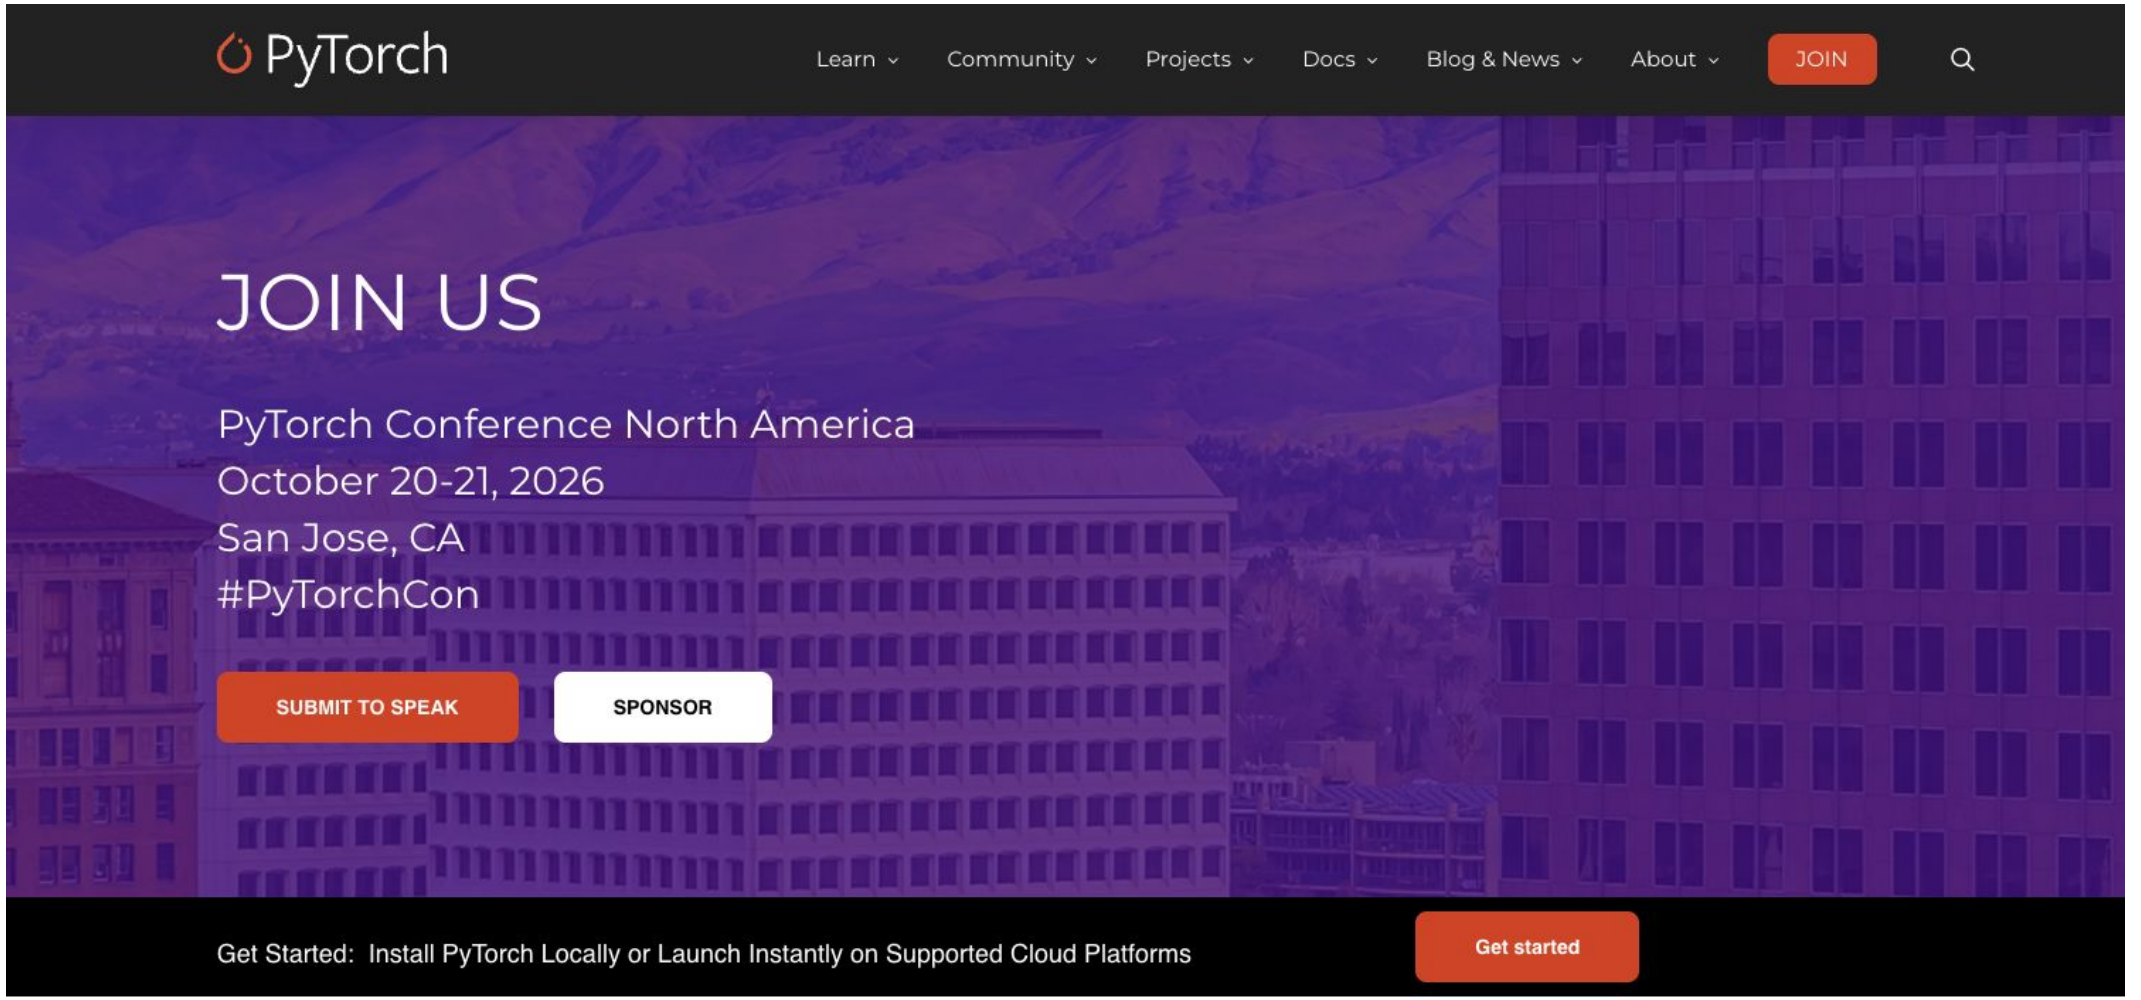
- January 13, 2026
- April 14, 2026
- July 14, 2026
- November 10, 2026 - updated



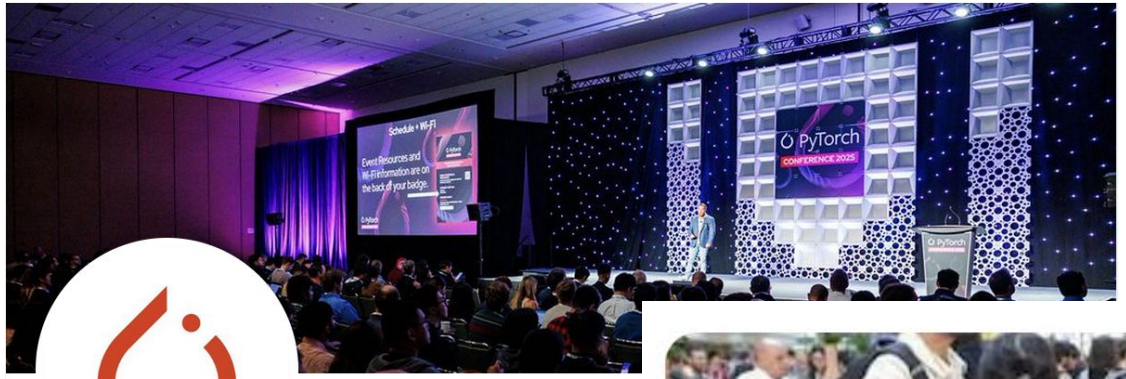
Events

Check out the
PyTorch Events Calendar

Marketing



Visit the PyTorch website: https://pytorch.org/



PyTorch
@PyTorch

Tensors and neural networks in Python with strc
PyTorch is an open source project at the Linux F

pytorch.org Joined September 2016 >

86 Following **496.5K** Followers



What PyTorch Conference Europe 2026 Was Really Like – Official PyTorchCon EU ...

PyTorch
321,544 followers
1w · Edited

Using PyTorch allowed LinkedIn to enable production-grade optimization at previously infeasible scales.

In our latest PyTorch Foundation case study, the engineering team at LinkedIn shares how they re-architected their extreme scale Linear Programming Solver (DuaLip GPU). Built using PyTorch to enable multi-GPU computations and parallelism, this solver serves as a core building block for massive matching problems that power ranking, personalization, and recommendation systems.

To maximize parallelism, the team tailored GPU execution techniques to sparse matching constraints. They developed constraint-aligned sparse layouts, batched projection kernels, and a distributed design focused on dual variables. They also updated the underlying ridge-regularized dual ascent method with Jacobi-style row normalization, primal scaling, and a regularization parameter continuation scheme.

When tested against extreme scale matching workloads, the GPU implementation achieved a 10x wall-clock speedup over the previous distributed CPU DuaLip solver while maintaining strict convergence guarantees. This project represents an excellent intersection of ML systems, mathematical optimization, and machine learning.

Read the complete case study here 👉 https://lnkd.in/eE67sxcB

Aida Rahmattalabi Sanjana Garg Zhipeng Wang, PhD Shenyinying Ruby Tu Yuan G. Gregory Dexter



Follow PyTorch on social media: X, LinkedIn, Facebook, Bluesky, YouTube



Subscribe to the PyTorch Newsletter to get the latest updates, event info, and community news in your inbox each month → pytorch.org/newsletter

PyTorch Trademark Policy

PyTorch® is a registered trademark. Per the Linux Foundation Trademark Usage Guidelines:

Trademarks are used to provide assurance to the community of consistency with respect to the quality of products or services with which the mark is associated. This policy attempts to strike the proper balance between two competing interests:

1) the need of The Linux Foundation to ensure that its trademarks remain reliable indicators of the qualities that they have been created to preserve; and,

2) the need of The Linux Foundation to ensure that community members are able to discuss the projects with which The Linux Foundation is associated and to accurately describe the relationship between The Linux Foundation and the products and services offered by others.

Questions? Email trademark@pytorch.org

- **Don't** use a trademark as a verb or noun - use only as an adjective followed by the generic name/noun.
- **Don't** use in the plural or possessive form.
- **Don't** alter or amend a trademark in any way, including combining with another mark, hyphenating, abbreviating, or displaying in parts.
- **Don't** use it as your domain or as part of your domain name.
- **Don't** use it as part of your product name.
- **Don't** incorporate it into your company's logos or designs.
- **Do** make fair use of word marks to make true factual statements, as long as you acknowledge the owner of the trademark.
- **Do** use the full form of any trademark in the first reference, after which you may then use any abbreviated or short form references.

For the full trademark guidance, please review https://www.linuxfoundation.org/legal/trademark-usage and https://pytorch.org/wp-content/uploads/2025/09/pytorch_brand_guide_091925a.pdf

Legal Notices

The Linux Foundation, The Linux Foundation logos, and other marks that may be used herein are owned by The Linux Foundation or its affiliated entities, and are subject to The Linux Foundation's Trademark Usage Policy at https://www.linuxfoundation.org/trademark-usage, as may be modified from time to time.

Linux is a registered trademark of Linus Torvalds. Please see the Linux Mark Institute's trademark usage page at https://lmi.linuxfoundation.org for details regarding use of this trademark.

Some marks that may be used herein are owned by projects operating as separately incorporated entities managed by The Linux Foundation, and have their own trademarks, policies and usage guidelines.

TWITTER, TWEET, RETWEET and the Twitter logo are trademarks of Twitter, Inc. or its affiliates.

Facebook and the "f" logo are trademarks of Facebook or its affiliates.

LinkedIn, the LinkedIn logo, the IN logo and InMail are registered trademarks or trademarks of LinkedIn Corporation and its affiliates in the United States and/or other countries.

YouTube and the YouTube icon are trademarks of YouTube or its affiliates.

All other trademarks are the property of their respective owners. Use of such marks herein does not represent affiliation with or authorization, sponsorship or approval by such owners unless otherwise expressly specified.

The Linux Foundation is subject to other policies, including without limitation its Privacy Policy at https://www.linuxfoundation.org/privacy and its Antitrust Policy at https://www.linuxfoundation.org/antitrust-policy. each as may be modified from time to time. More information about The Linux Foundation's policies is available at https://www.linuxfoundation.org.

Please email legal@linuxfoundation.org with any questions about The Linux Foundation's policies or the notices set forth on this slide.