PyTorch Foundation
Technical Advisory Council Meeting

2026 August 11

Reminder for meeting participants

- This meeting is being recorded
- Please remember to be on mute upon entry unless you would like to speak
- All zoom chat messages will become public
- Please ensure your name and affiliation are showing correctly

Questions? Email support@pytorch.org

Antitrust Policy Notice

Linux Foundation meetings involve participation by industry competitors, and it is the intention of the Linux Foundation to conduct all of its activities in accordance with applicable antitrust and competition laws. It is therefore extremely important that attendees adhere to meeting agendas, and be aware of, and not participate in, any activities that are prohibited under applicable US state, federal or foreign antitrust and competition laws.

Examples of types of actions that are prohibited at Linux Foundation meetings and in connection with Linux Foundation activities are described in the Linux Foundation Antitrust Policy available at http://www.linuxfoundation.org/antitrust-policy. If you have questions about these matters, please contact your company counsel, or if you are a member of the Linux Foundation, feel free to contact Andrew Updegrove of the firm of Gesmer Updegrove LLP, which provides legal counsel to the Linux Foundation.

Agenda

- Administrative – 5 min
- Foundation Updates – 15 min
- Foundation-Hosted Project Application: ONNX – 15 min
- Working Group Updates – 10 min
- Call for Topics/Open Discussion – 5 min

Meeting Schedule

2026

- ~~January 13~~
- ~~February 10~~
- ~~March 10~~
- ~~April 14~~
- ~~May 12~~
- ~~June 9~~
- ~~July 14~~
- August 11
- September 15
- October 13
- November 10
- December 8

All meetings start at 7am PT/4pm CET unless otherwise indicated.

Moved 1 week to avoid conflict with PyTorch Conference China.

Upcoming TAC Meetups

PyTorch Conference China

TAC Lunch Meetup (Members Only) – **RSVP to Calendar Invite by August 20th!**

Tuesday September 8 | 12:15pm – 1:45pm CST

Register to Attend Event

PyTorch Conference North America

TAC Meetup (Members Only) – **RSVP to Calendar Invite by October 5th!**

Wednesday October 21 | 12:45pm – 2:15pm PT

Register to Attend Event

PyTorch Conference Community Night (*Sponsored by Huawei*) – Monday, September 7 @ 17:30 - 20:30
Puzhixing, F4, Pujiang Hall – First come, first serve, limited capacity!

TAC Representatives - Member Companies

Welcome Gerred, Joseph, and Naigang!

Company/Project	Primary Voting Representative	Alternate Voting Representative
Alibaba Cloud	Tao Ma	Sanhong Li
Amazon Web Services	Shauheen Zahirazami	Sundar Ranganathan
AMD	Jeff Daily	Niles Burbank
ARM	Milos Puzovic (TAC Vice Chair)	Andrew Wafaa
Google Cloud	Claudio Basile	Michael Voznesensky
Huawei	Yikun Jiang	Fengchun Hua
IBM	Joseph Groenenboom	Naigang Wang
Intel	Eikan Wang	Ashok Emani
Lightning AI	Thomas Viehmann (TAC Chair)	
Meta	Gregory Chanan	Hamid Shojanazeri
Nscale	Gerred Dillon	
NVIDIA	Piotr Bialecki	Michael Ruberry
Snowflake	Mert Hidayetoglu	
Shopify	Shuying Sun	
Gold Members	Weizhu Chen	

TAC Representatives - Projects

Company/Project	Primary Voting Representative	Alternate Voting Representative
PyTorch Project	Alban Desmaison	
DeepSpeed	Tunji Ruwase	Minjia Zhang
vLLM	Simon Mo	Kaichao You
Ray	Robert Nishihara	
Helion	Jongsok “James” Choi	
Safetensors	Luc Georges	Lysandre Debut

Foundation Updates
Mark Collier

Time allocation: 15 minutes

2026 Goals

Evolve into a multi-project foundation

Future ready CI infra approach

Connect Community IRL

Project Lifecycle Policy - Status Update

The PyTorch Project Lifecycle Policy has been in the process of being updated, however LF legal has recommended considering simpler approaches.

Link to current draft: PyTorch Project Lifecycle Policy (version 3.4)

Next steps: review potential shorter/standard template options linked below

- Draft based on ASWF's policy - link
- Draft based on standard LF lifecycle template - link

Additional options: hold one-off working meeting to complete draft doc, review asynchronously, ? (discussion)

PyTorch Foundation Ambassador Program

Launched in 2025, the PyTorch Foundation Ambassador Program supports community leaders who expand the global reach of PyTorch Foundation-hosted projects through events, education, and advocacy.

Program Highlights

- 31 active ambassadors across 20 countries
- October 2025–March 2026: 54 community activities delivered globally, including 36 events, 9 blogs/articles, 6 trainings/workshops, and 3 videos.
- New regional communities launched, including PyTorch Nepal, PyTorch East Africa, PyTorch Nigeria, PyTorch Michigan, and PyTorch Sri Lanka
- The 2026 Ambassador recruitment cycle aims to recruit approximately 20 new ambassadors to expand regional representation and strengthen engagement across PyTorch Foundation-hosted projects.

2026 Program Key Dates

- May 18, 2026 – Ambassador nominations open
- June 19, 2026 – Ambassador nominations close
- July 17, 2026 – Reviewer Interest Form deadline
- August 7- 17 2026 – Reviewer evaluation period
- September 2026 - Application Notifications

Foundation-Hosted Project Application: ONNX
Andreas Fehlner

Time allocation: 15 minutes

Working Group Updates

Time allocation: 10 minutes

CI Working Group Update
Thanh Ha

CI Working Group Update

- AWS Costs Update
  - Costs in July: 747k
    - We allowed higher spend in July to make up for underspend in May/June
  - AWS Credits remaining: 1.04m
  - 25% of jobs allocated to Foundation account (Excluding Inductor Jobs)
- GitHub OSDC (ARC) CI Updates
  - Continuing to optimize costs and improve cluster sizing
  - Improved OSDC runner utilization by roughly 5%
  - Cleaned up legacy MultiCloud WG ARC dev cluster as OSDC has replaced it

Costs To PyTorch Foundation by Architecture July 2026

Cost To PyTorch Foundation By Architecture	
Vendor	**Total**
x86_64 (AMD)	$296,870.80
CUDA	$220,526.30
N/A	$132,838.14
ARM64 (AWS Graviton)	$47,958.79
x86_64 (Generic)	$42,944.37
Windows	$4,447.61
x86_64 (Intel)	$1,024.81
XPU	$299.31
ARM64 (Unknown)	$0.00
s390x	$0.00
x86_64 (Unknown)	$0.00
ROCM	$0.00
TPU	$0.00
macOS	$0.00
	$746,910.13

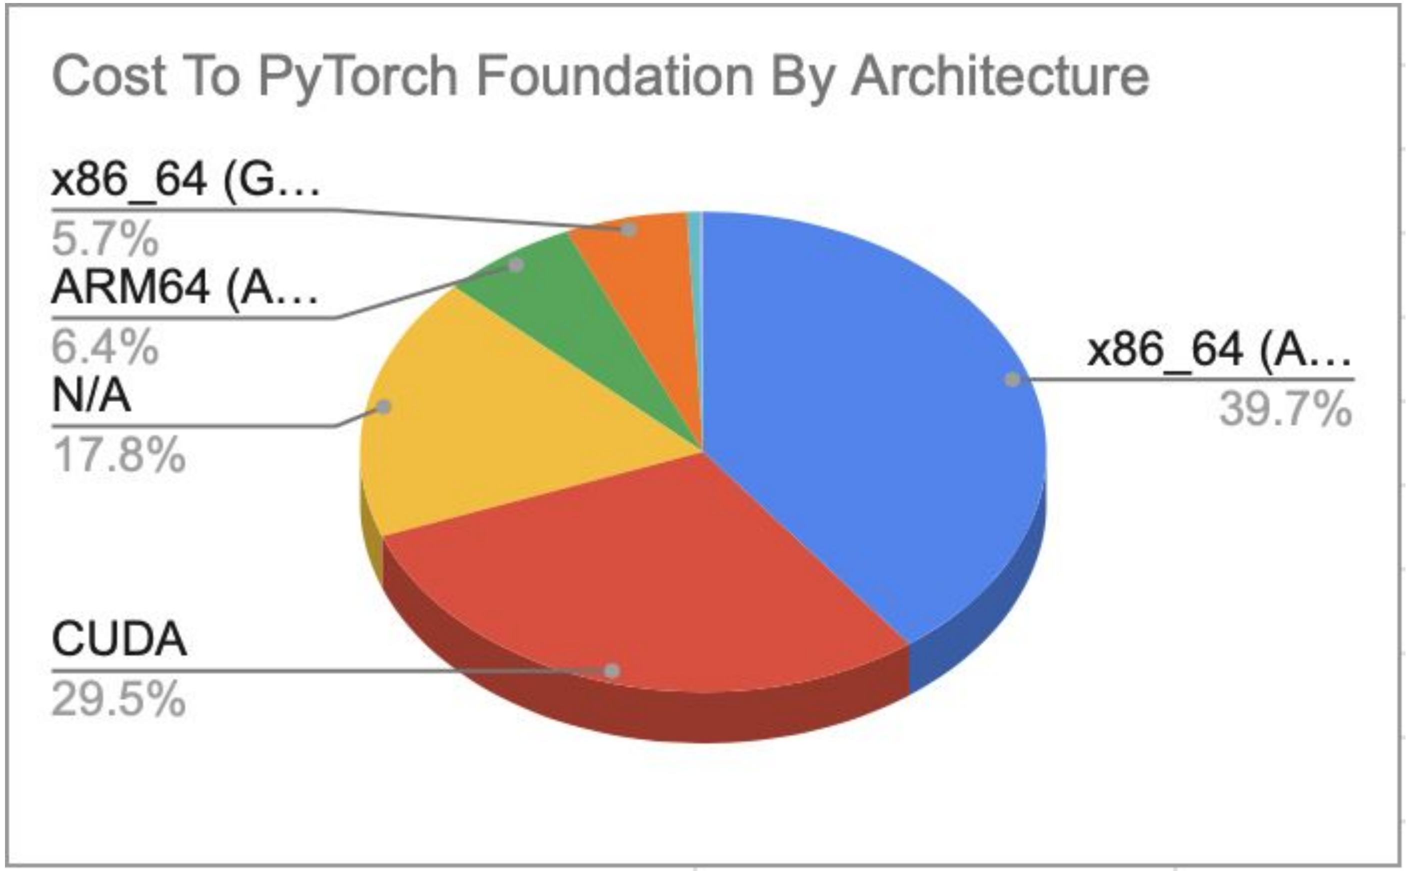

* Data from PT Foundation AWS Account. This does not include self-hosted runner costs by member provided runners.

Combined Hardware Runtime by Architecture 12 month trend (1/2)

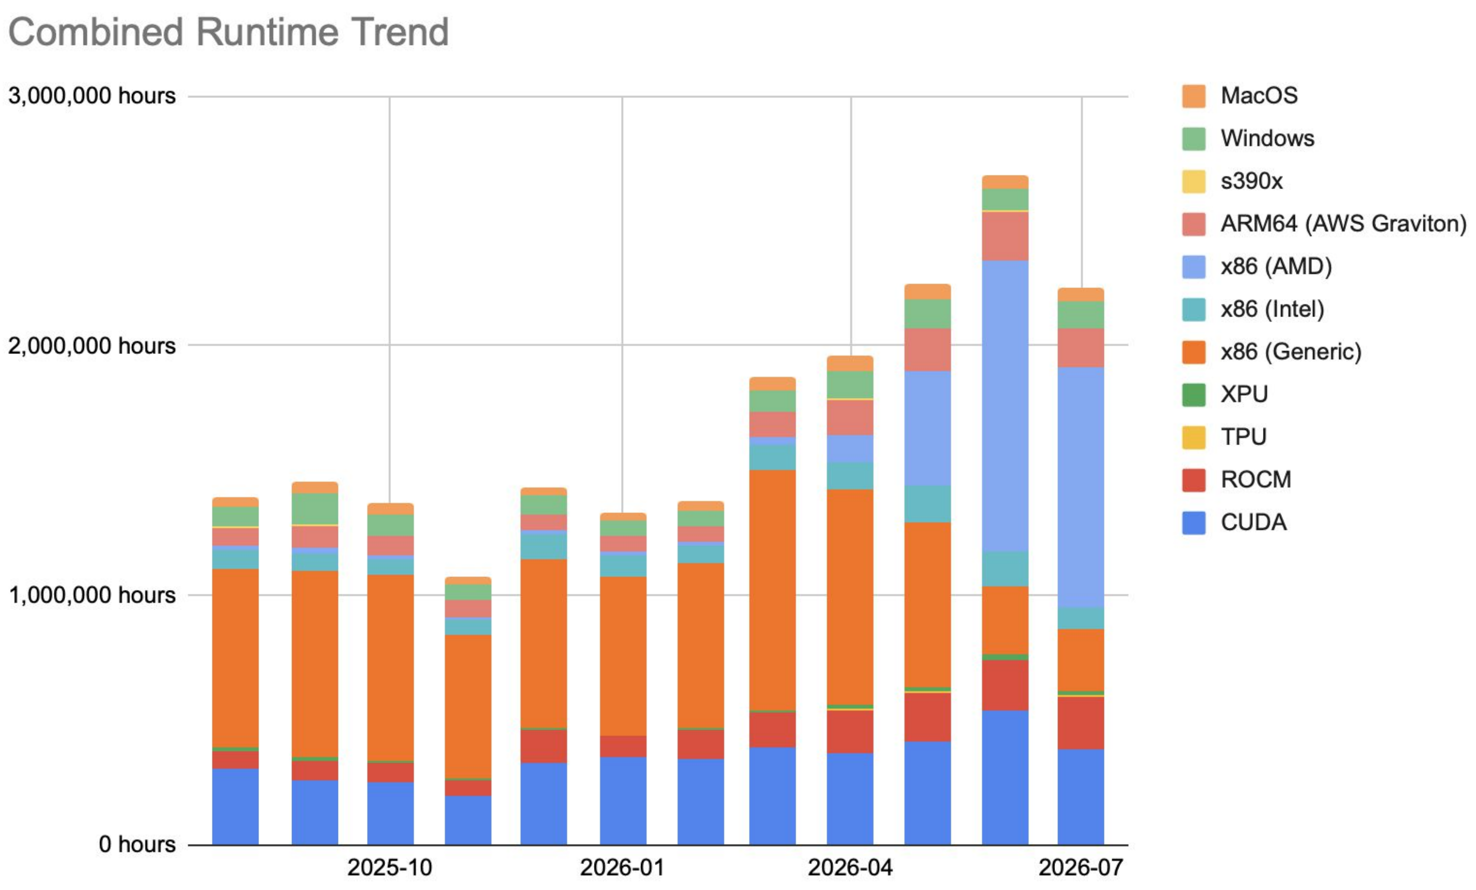

* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners.

Combined Hardware Runtime by Architecture 12 month trend (2/2)

Vendor	2025-08	2025-09	2025-10	2025-11	2025-12	2026-01	2026-02	2026-03	2026-04	2026-05	2026-06	2026-07
ARM64 (AWS Graviton)	68,099 hours	90,597 hours	74,757 hours	65,201 hours	66,970 hours	60,192 hours	59,735 hours	98,649 hours	142,832 hours	166,475 hours	200,398 hours	158,681 hours
ARM64 (Unknown)	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	7 hours	0 hours	0 hours	0 hours	0 hours	0 hours
s390x	4,502 hours	6,245 hours	2,113 hours	1,964 hours	2,090 hours	2,221 hours	3,157 hours	2,632 hours	8,072 hours	5,650 hours	2,322 hours	3,451 hours
x86 (AMD)	18,463 hours	16,791 hours	16,212 hours	12,002 hours	15,715 hours	13,841 hours	14,649 hours	34,247 hours	111,862 hours	460,903 hours	1,163,444 hours	958,747 hours
x86 (Intel)	76,809 hours	69,760 hours	63,913 hours	60,868 hours	97,881 hours	90,876 hours	70,572 hours	97,444 hours	104,909 hours	146,231 hours	138,474 hours	89,339 hours
x86 (Generic)	718,661 hours	751,172 hours	740,671 hours	573,826 hours	675,526 hours	632,118 hours	661,004 hours	965,287 hours	866,451 hours	664,920 hours	274,431 hours	244,337 hours
x86 (Unknown)	5 hours	98 hours	101 hours	103 hours	114 hours	50 hours	66 hours	675 hours	557 hours	649 hours	711 hours	744 hours
CUDA	306,333 hours	257,586 hours	250,679 hours	199,004 hours	331,416 hours	350,926 hours	344,301 hours	387,160 hours	364,590 hours	410,558 hours	537,420 hours	379,480 hours
ROCM	66,102 hours	76,163 hours	79,219 hours	59,104 hours	129,254 hours	83,308 hours	113,864 hours	145,289 hours	175,847 hours	200,134 hours	204,177 hours	216,072 hours
TPU	0 hours	0 hours	0 hours	0 hours	73 hours	125 hours	452 hours	877 hours	1,663 hours	1,194 hours	810 hours	870 hours
XPU	14,110 hours	16,320 hours	9,495 hours	9,168 hours	8,062 hours	6,380 hours	11,156 hours	7,254 hours	16,958 hours	17,442 hours	19,649 hours	21,643 hours
MacOS	35,655 hours	46,812 hours	46,614 hours	29,719 hours	33,700 hours	32,748 hours	38,253 hours	52,769 hours	57,907 hours	62,271 hours	51,256 hours	53,192 hours
Windows	85,023 hours	126,600 hours	86,312 hours	63,557 hours	70,322 hours	57,758 hours	61,080 hours	85,558 hours	108,340 hours	112,295 hours	88,996 hours	104,749 hours
N/A	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours	0 hours
Total	1,325,662 hours	1,367,545 hours	1,295,329 hours	1,009,315 hours	1,364,152 hours	1,270,350 hours	1,318,561 hours	1,779,193 hours	1,817,156 hours	2,082,247 hours	2,481,690 hours	2,072,624 hours

* Data from PT Foundation AWS Account + Estimates from PyTorch HUD. This does not include self-hosted runner costs by member provided runners.

Ecosystem Working Group
Chris Hoge/Eikan Wang

Ecosystem Working Group Updates

Accepted 🎉:

- 🎉 TokenSpeed – Lightseek PyTorch-based LLM inference engine optimized for agentic workloads – needs
- 🎉 FiftyOne – Open-source tool for building high-quality datasets and computer vision models – ready for vote
- 🎉 VisualTorch – PyTorch model visualization library
- 🎉 SMG – Lightseek Shepherd Model Gateway – voting underway

Active Applications 🗳️:

- 🗳️ TorchSurv – "Deep survival analysis made easy" – in voting
- 📥 NeuralDBG – NeuralDBG is a causal diagnostic engine for PyTorch training – needs intake review

Under Review 🔍:

- ⚠️ torchao – easy to use quantization library for native PyTorch – **part of PyTorch GitHub, missing maintainers**
- ⚠️ Zeus – library for measuring the energy consumption of Deep Learning workloads – in remediation

Upcoming Workstreams 🛠️:

- Writing new WG charter
- Finalizing survey
- Developing skills and guidelines for consistent and repeatable review across WG members.
- Preparing talk for PyTorch Conference
- Agent accessibility on landscape

Monthly meeting on last Wednesday
at 9am PT/12pm ET – link
https://github.com/pytorch-fdn/ecosystem
Slack: #pytorch-ecosystem-working_group

Accelerator Integration Working Group

Zesheng Zong

Accelerator Integration Working Group Updates

Additional Compute Platforms Page

- The page is officially launched. (get-started/additional-platforms/)
- Ascend NPU and Intel XPU have passed the voting-member review, and their introduction content is now live.

Reference

- Check [Additional Compute Platforms Admission Process]()
- Submit applications in the [additional-compute-platforms]() repo
- WG Charter: [governance/CHARTER.md]()

CTA: TAC to review/approve charter via LFX vote

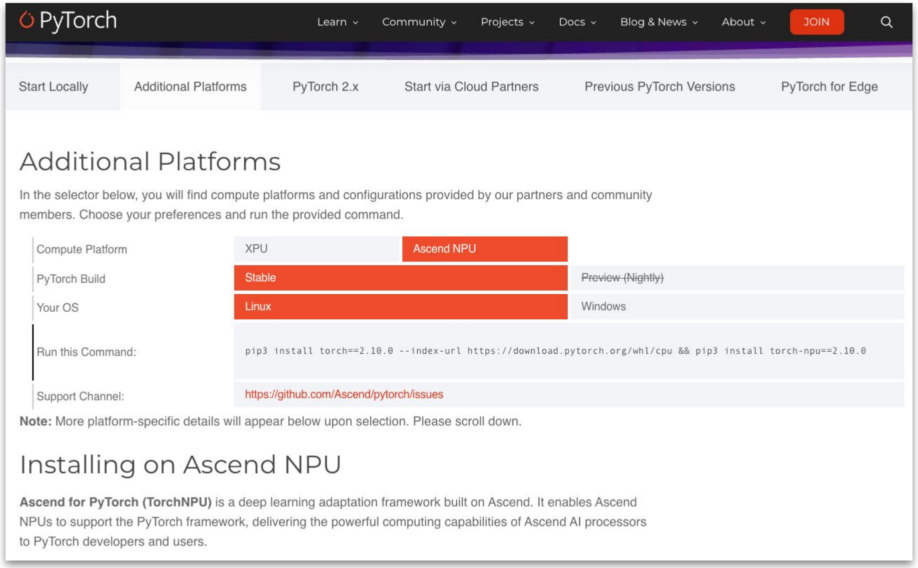

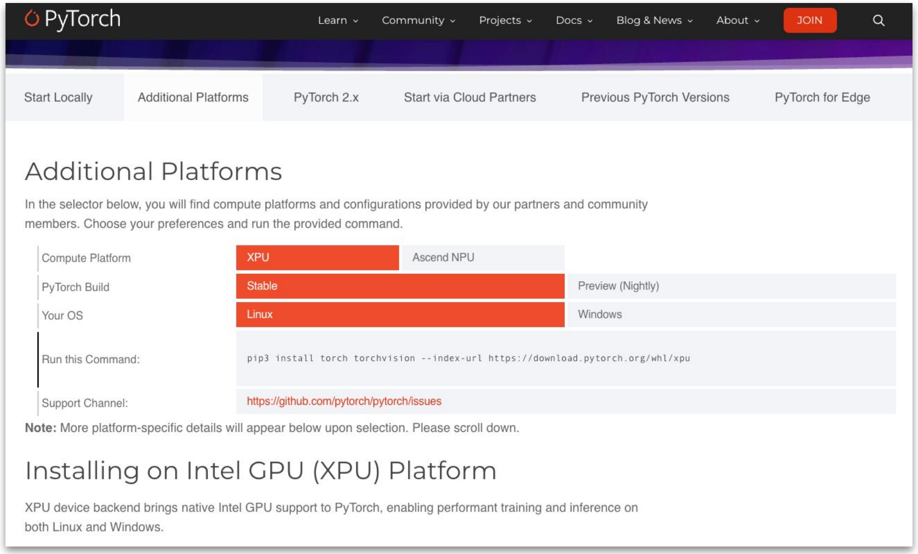

Accelerator Integration Working Group Updates

Cross-Repository CI Relay (RFC)

- HUD improvements shipped– Metrics page, per-repo dashboards, Healthy/Degraded status, L3/L4 grid filtering, and idle crash fix. (#8318,#8319,#8330,#8341,#8343,#8366)
- Dr.CI and oncall bot integration– CRCR failures now surface in Dr.CI comments with job names, and cross-fork PR resolution added. (#8365,#8386)
- Nightly/Periodic CI in progress– Callback handler, SHA validator, and action inputs merged. ClickHouse schema and HUD nightly tab under development. (rfcs#98,#8302,#8303,#8304,#8353,#8377)
- L4 blocking workflow in progress– Blocking check runs on upstream PRs for L4 repos. (crcr-test#18), and the testing design for L4 is under the discussion.
- Dispatch filtering live– Downstream CI now triggers only on merged PRs orciflow/crcr-labeled PRs, preventing unnecessary runs. (crcr-test#14)
- OOT → CRCR rename complete– Trigger labels, job names, and RFC terminology fully migrated to CRCR branding. (rfcs#100,crcr-test#19)

Test Refactoring

- 273/1,023 test files have been refactored, with 130+ PRs under initial review and 7 PRs under final review.
- Key Progress: PR #190173 is well underway to add linter for test case refactoring.

Security Working Group
Nikita Shulga

Security WG Updates

- 18 issues reported, 15 closed, 1 in draft
  - Allowing `@pytorchbot lint` to all contributors is indeed insecure, restricted to users with write permissions

- Incoming advisories are triaged every other week at 8:30AM PST
  - Meetings are not open to public by design
  - If you have some security background and want to join, post in `#tac-security-wg`

OSPO & Academic Outreach Working Group

Ana Jimenez/George Chellapa/Sumantro Mukherjee

OSPO & Academic Outreach WG Updates

Repo: https://github.com/pytorch-fdn/wg-ospo-and-academic-outreach

WG Chairs: Sumantro Mukherjee (Red Hat), George Chellapa (NVIDIA), Ana Jiménez

Updates:

- Curriculum Development Workstream charter is live!
- Academic Survey is being finalized in collaboration with CURIOUS, WG leads and university academic reps.
- Plans to host a OSPO & Academic Outreach meetup/event as a precursor to PyTorch Conference NA.
- Pilot blog highlighting TinyTorch being drafted in collaboration with Harvard University.
- Style guide for Academic OSPO blog posts is live!
- OSPO & Academic WG will closely work with Marketing Teams for content publication and visibility.
- Invitations to academic professors/students to participate in WG to be launched for increased participation in WG

Open Discussion

Open Discussion

- Any member can take the floor

Next Meeting

Next regular TAC meeting

Tuesday, September 15, 2026

7am PDT / 10am EDT / 4pm CEST

Topics: tbd

Please ensure you are registered for the meeting series and that meetings are reflected in your calendar via LFX Meetings

Have an agenda topic idea? Post topics to the TAC Agenda Github Project

Helpful Links

Links

Foundation GitHub:
https://github.com/pytorch-fdn/tac

Mailing List:
https://lists.pytorch.org/g/tac

Past Meeting Minutes and Decks:
https://github.com/pytorch-fdn/tac/tree/main/minutes

LFX Individual Dashboard:
https://openprofile.dev

Submit an Ecosystem project:
https://github.com/pytorch-fdn/ecosystem

Submit a Blog:
https://form.asana.com/?k=mPrsX9EcZ0BoGnqz543C1A&d=9283783873717

Points of Contact

Mark Collier (Executive Director)

Naomi Washington (Senior Program Manager)

Jennifer Bly (Marketing Director)

Regina Nkenchor (Program Manager)

Michelle Roth (Program Manager)

Questions? Email support@pytorch.org

Appendix

Foundation-Hosted Project Review Schedule

- January 13, 2026
- April 14, 2026
- July 14, 2026
- November 10, 2026 - updated

Events

Check out the

PyTorch Events Calendar

Marketing

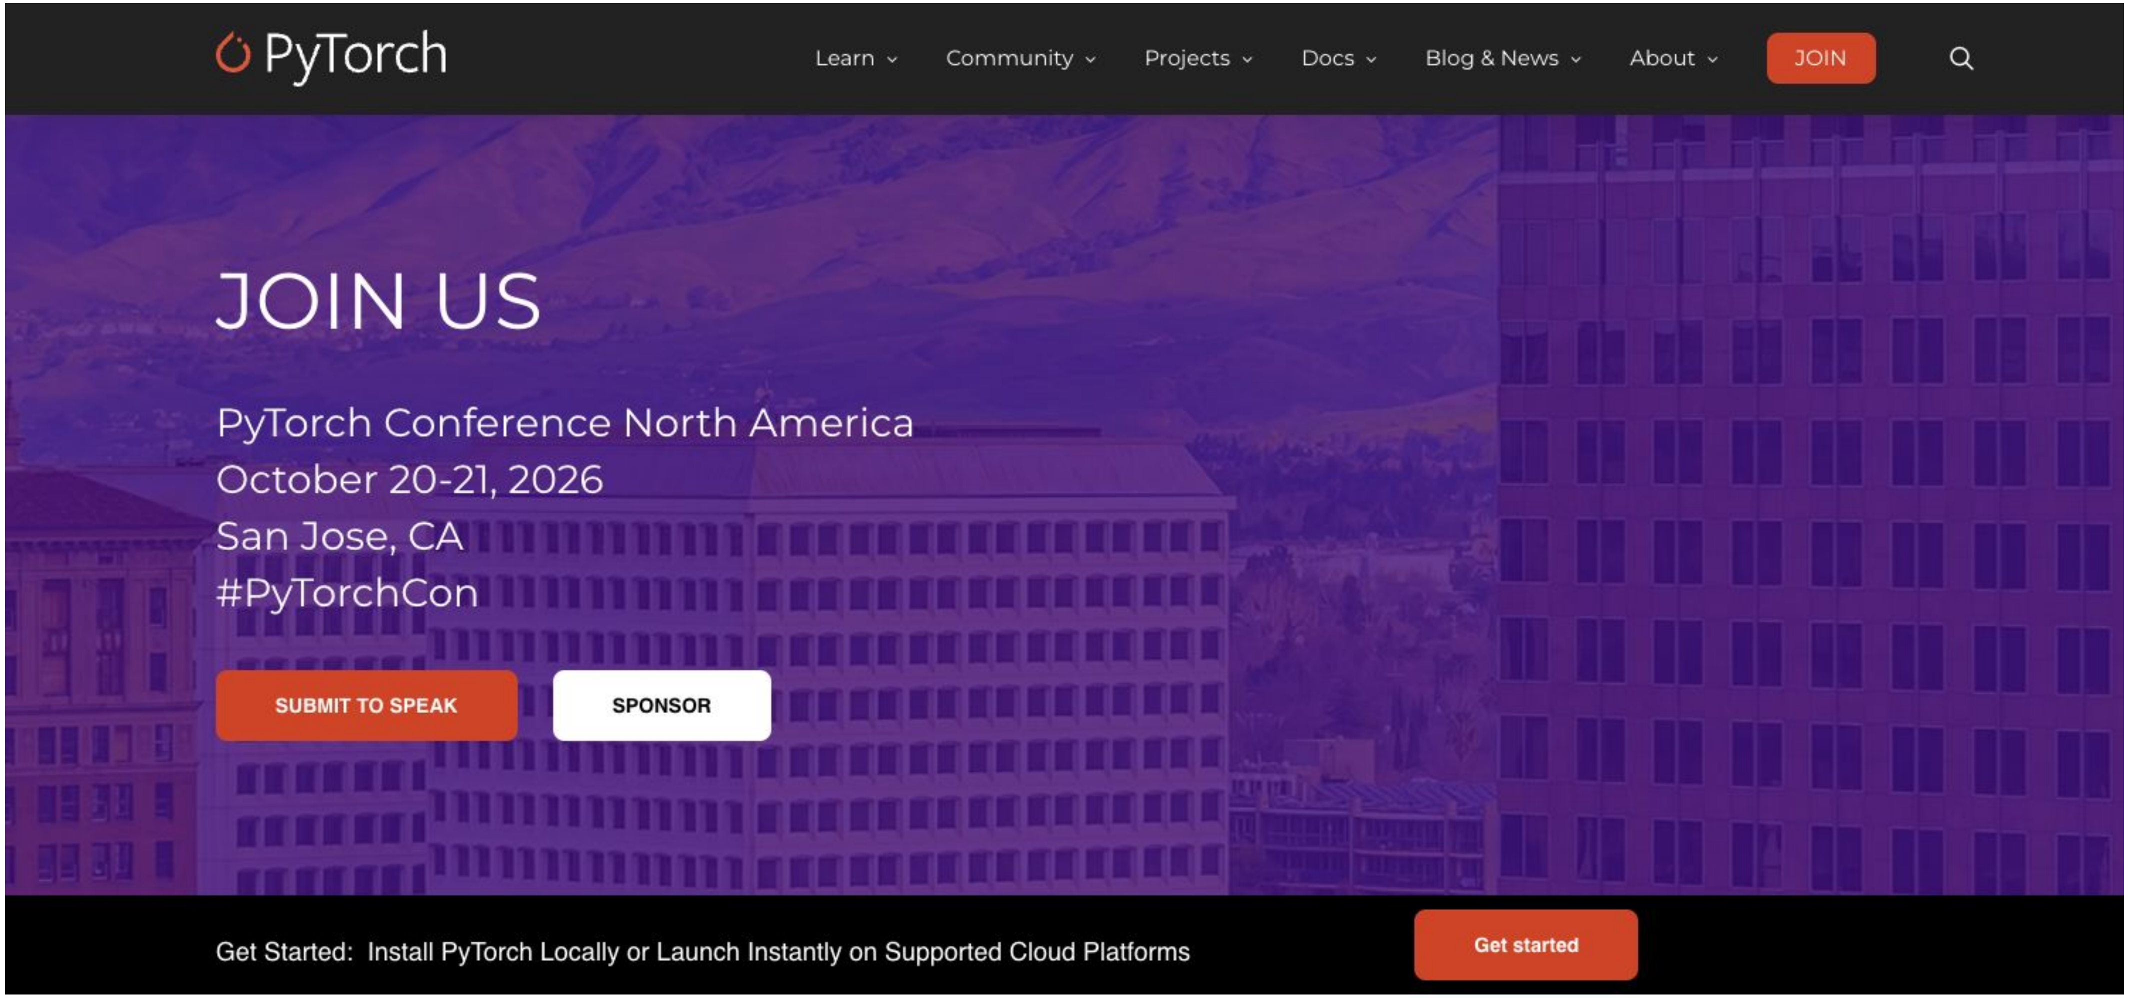

Visit the PyTorch website: https://pytorch.org/

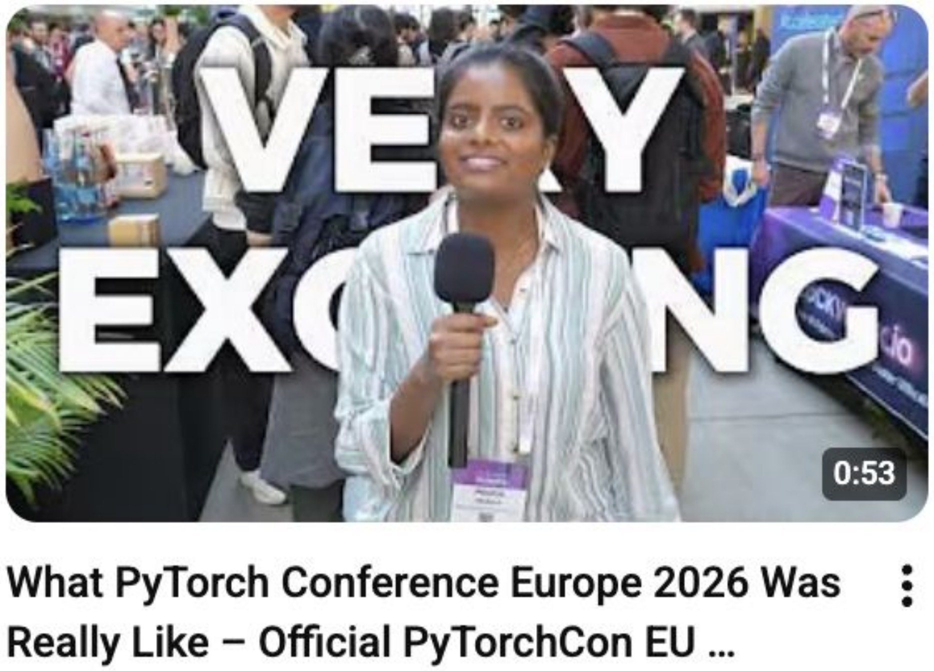

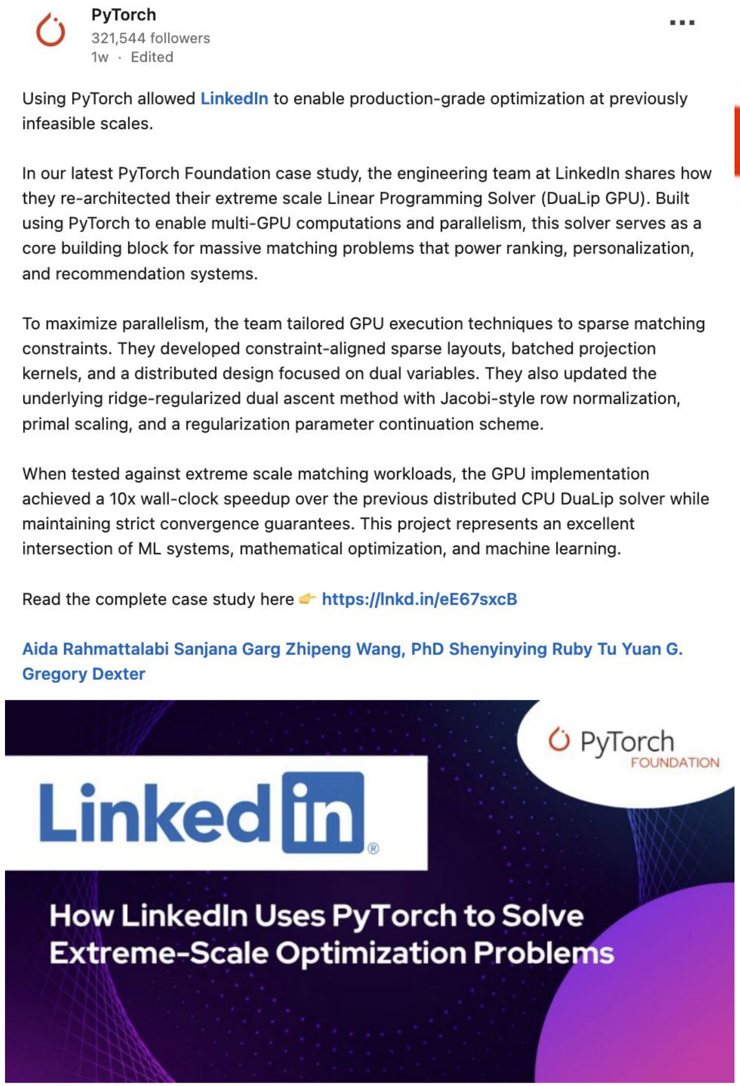

Follow PyTorch on social media: X, LinkedIn, Facebook, Bluesky, YouTube

Subscribe to the PyTorch Newsletter to get the latest updates, event info, and community news in your inbox each month → pytorch.org/newsletter

PyTorch Trademark Policy

PyTorch® is a registered trademark. Per the Linux Foundation Trademark Usage Guidelines:

Trademarks are used to provide assurance to the community of consistency with respect to the quality of products or services with which the mark is associated. This policy attempts to strike the proper balance between two competing interests:

1) the need of The Linux Foundation to ensure that its trademarks remain reliable indicators of the qualities that they have been created to preserve; and,

2) the need of The Linux Foundation to ensure that community members are able to discuss the projects with which The Linux Foundation is associated and to accurately describe the relationship between The Linux Foundation and the products and services offered by others.

Questions? Email trademark@pytorch.org

- **Don't** use a trademark as a verb or noun - use only as an adjective followed by the generic name/noun.
- **Don't** use in the plural or possessive form.
- **Don't** alter or amend a trademark in any way, including combining with another mark, hyphenating, abbreviating, or displaying in parts.
- **Don't** use it as your domain or as part of your domain name.
- **Don't** use it as part of your product name.
- **Don't** incorporate it into your company's logos or designs.
- **Do** make fair use of word marks to make true factual statements, as long as you acknowledge the owner of the trademark.
- **Do** use the full form of any trademark in the first reference, after which you may then use any abbreviated or short form references.

For the full trademark guidance, please review https://www.linuxfoundation.org/legal/trademark-usage and https://pytorch.org/wp-content/uploads/2025/09/pytorch_brand_guide_091925a.pdf

Legal Notices